

```
# gerardo Herrera... random forest (25 arboles) con 28k instancias de normal y recovering y 24
```

```
from google.colab import drive
drive.mount('/content/drive')
```

```
Mounted at /content/drive
```

```
import numpy as np
import pandas as pd
import os
```

```
import matplotlib.pyplot as plt
%matplotlib inline
from tqdm import tqdm_notebook
from sklearn.preprocessing import StandardScaler
from sklearn.metrics import mean_absolute_error
pd.options.display.precision = 15
```

```
import time
# Libraries
import numpy as np
import pandas as pd
pd.set_option('max_columns', None)
import matplotlib.pyplot as plt
import seaborn as sns
%matplotlib inline
```

```
!pip install lightgbm
!pip install catboost
```

```
import datetime
import lightgbm as lgb
from scipy import stats
from sklearn.model_selection import train_test_split, StratifiedKFold, KFold, cross_val_score
from sklearn.preprocessing import StandardScaler
import os
import lightgbm as lgb
from sklearn.metrics import accuracy_score, roc_auc_score
from sklearn import metrics
from sklearn import linear_model
from tqdm import tqdm_notebook
from catboost import CatBoostClassifier
```

```
↳ Requirement already satisfied: lightgbm in /usr/local/lib/python3.6/dist-packages (2.2.1)
Requirement already satisfied: scikit-learn in /usr/local/lib/python3.6/dist-packages (0.22.2)
Requirement already satisfied: scipy in /usr/local/lib/python3.6/dist-packages (from lightgbm) (1.4.1)
Requirement already satisfied: numpy in /usr/local/lib/python3.6/dist-packages (from lightgbm) (1.19.5)
```

```
Requirement already satisfied: joblib>=0.11 in /usr/local/lib/python3.6/dist-packages (1
Collecting catboost
  Downloading https://files.pythonhosted.org/packages/52/39/128fff65072c8327371e3c594f3c
|████████████████████████████████████████████████████████████████████████████████| 66.2MB 56kB/s
Requirement already satisfied: six in /usr/local/lib/python3.6/dist-packages (from catb
Requirement already satisfied: scipy in /usr/local/lib/python3.6/dist-packages (from cat
Requirement already satisfied: numpy>=1.16.0 in /usr/local/lib/python3.6/dist-packages (
Requirement already satisfied: graphviz in /usr/local/lib/python3.6/dist-packages (from
Requirement already satisfied: pandas>=0.24.0 in /usr/local/lib/python3.6/dist-packages
Requirement already satisfied: matplotlib in /usr/local/lib/python3.6/dist-packages (fr
Requirement already satisfied: plotly in /usr/local/lib/python3.6/dist-packages (from c
Requirement already satisfied: pytz>=2017.2 in /usr/local/lib/python3.6/dist-packages (
Requirement already satisfied: python-dateutil>=2.7.3 in /usr/local/lib/python3.6/dist-p
Requirement already satisfied: kiwisolver>=1.0.1 in /usr/local/lib/python3.6/dist-packag
Requirement already satisfied: cycler>=0.10 in /usr/local/lib/python3.6/dist-packages (
Requirement already satisfied: pyparsing!=2.0.4,!=2.1.2,!=2.1.6,>=2.0.1 in /usr/local/li
Requirement already satisfied: retrying>=1.3.3 in /usr/local/lib/python3.6/dist-packages
Installing collected packages: catboost
Successfully installed catboost-0.24.2
```



```
# sensor = pd.read_csv('../input/sensor.csv')
# sensor = pd.read_csv('../input/vombas/sensor_procesado.csv')
#sensor = pd.read_csv('dataset_sensor_procesado.csv')
#sensor = pd.read_csv('../input/bombas-sensores-conocidos/sensor2.csv')
#sensor = pd.read_csv('../input/28k-s24-balan-vombas/sensor2-ordenado_status_sin_broken_balan')
#sensor.drop(['Unnamed: 0'], axis=1, inplace=True)

sensor = pd.read_csv('/content/drive/My Drive/datasets/sensor2-ordenado_status_sin_broken_bal

sensor.head()
```


	Unnamed: 0	sensor_00	sensor_01	sensor_02	sensor_03
0	0	2.465394	47.0920100000000002	53.2117999999999997	46.3107600000000002
1	1	2.465394	47.0920100000000002	53.2117999999999997	46.3107600000000002
2	2	2.444734	47.3524299999999998	53.2117999999999997	46.3975700000000002
3	3	2.460474	47.0920100000000002	53.1683999999999998	46.397567749023402
4	4	2.445718	47.1354100000000000	53.2117999999999997	46.397567749023402
5	5	2.453588	47.0920100000000002	53.1683999999999998	46.397567749023402
6	6	2.455556	47.0486099999999996	53.168399810790994	46.397567749023402
7	7	2.449653	47.1354100000000000	53.168399810790994	46.397567749023402
8	8	2.463426	47.0920100000000002	53.168399810790994	46.397567749023402
9	9	2.445718	47.1788200000000002	53.1683999999999998	46.397567749023402
10	10	2.464410	47.4826400000000004	53125.0000000000000000	46.397567749023402
11	11	2.444734	47.9166600000000000	53.1683999999999998	46.397567749023402
12	12	2.460474	48.2638900000000004	53125.0000000000000000	46.397567749023402
13	13	2.448669	48.4375000000000000	53.1683999999999998	46.397567749023402
14	14	2.453588	48.5677099999999998	53.1683999999999998	46.397567749023402
15	15	2.455556	48.3941000000000002	53125.0000000000000000	46.3975700000000002
16	16	2.449653	48.3941000000000002	53.1683999999999998	46.3107600000000002
17	17	2.463426	48.4808999999999998	53.6892400000000012	46.310760498046896
18	18	2.445718	48.6111099999999996	53125.0000000000000000	46.310760498046896
19	19	2.464410	48.6111099999999996	53.1683999999999998	46.310760498046896

```
for col in sensor.columns[1:-1]:
```

```
    sensor[col] = sensor[col].fillna(sensor[col].mean())
```

```
..      ..      2.453588      49.0885400000000002      53.1683999999999998      46.267360687255895
```

```
# bosque aleatorio
```

```
23      23      2.453588      49.0885400000000002      53.1683999999999998      46.267360687255895
```

```
sensor.fillna(sensor.mean(), inplace=True)
```

```
sensor.head()
```

	Unnamed: 0	sensor_00	sensor_01	sensor_02	sensor_03	
0	0	2.465394	47.0920100000000002	53.2117999999999997	46.3107600000000002	6343
1	1	2.465394	47.0920100000000002	53.2117999999999997	46.3107600000000002	6343
2	2	2.444734	47.3524299999999998	53.2117999999999997	46.3975700000000002	6

```
print(sensor.shape)
```

```
(28002, 26)
```

```
# Encontrar características importantes en Scikit-learn
```

```
# from sklearn.ensemble import RandomForestClassifier
```

```
#Create a Gaussian Classifier
```

```
#clf=RandomForestClassifier(n_estimators=100)
```

```
#Train the model using the training sets y_pred=clf.predict(X_test)
```

```
#clf.fit(X_train,y_train)
```

```
# no correr
```

```
#import pandas as pd
```

```
#feature_imp = pd.Series(clf.feature_importances_,index=iris.feature_names).sort_values(ascen
```

```
#feature_imp = pd.Series(clf.feature_importances_,index=sensor.columns[19:27]).sort_values(as
```

```
#print(feature_imp)
```

```
#Visualización
```

```
#import matplotlib.pyplot as plt
```

```
#import seaborn as sns
```

```
##matplotlib inline
```

```
# Creating a bar plot
```

```
#sns.barplot(x=feature_imp, y=feature_imp.index)
```

```
# Add labels to your graph
```

```
#plt.xlabel('Feature Importance Score')
```

```
#plt.ylabel('Features')
```

```
#plt.title("Visualizing Important Features")
```

```
#plt.legend()
```

```
#plt.show()
```

```
X=sensor[['sensor_00', 'sensor_01', 'sensor_02', 'sensor_03','sensor_04', 'sensor_11', 'senso
```

```
#y=sensor['target'] # Labels
```

```
y=sensor['machine_status'] # Labels
```

```
# Split dataset into training set and test set
```

```
#X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3) # 70% training and
```

```

# Split dataset into training set and test set
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2) # 80% training and 20% test

from sklearn.ensemble import RandomForestClassifier

#Create a Random Forest Classifier
clf=RandomForestClassifier(n_estimators=25)

start = time.time()

#Train the model using the training sets y_pred=clf.predict(X_test)
clf.fit(X_train,y_train)

stop = time.time()
print(f"Training time: {stop - start}s")

y_pred=clf.predict(X_test)

#Import scikit-learn metrics module for accuracy calculation
from sklearn import metrics
# Model Accuracy, how often is the classifier correct?
print("Accuracy:",metrics.accuracy_score(y_test, y_pred))

#predicciones del item 17156 q es 1
clf.predict([[0.0,53.55902,52.77777,43.402774810790994,204.72509765625,3.7302410000000004,404

    Training time: 0.7204201221466064s
    Accuracy: 0.9998214604534904
    array([1])

#predicciones
clf.predict([[0.0,53.55902,0,0,0,0,0,0,0,0,0,0,0,0,0,0,0,0,0,0]])

    array([1])

#predicciones
clf.predict([[0.0,53.55902,52.77777,43.402774810790994,204.72509765625,3.7302410000000004,404

    array([1])

# Extract single tree
estimator = clf.estimators_[5]

#from sklearn.tree import export_graphviz
# Export as dot file
#export_graphviz(estimator, out_file='tree.dot',
#
#                 feature_names = ['sensor 00'. 'sensor 01'. 'sensor 02'. 'sensor 03'. 'sensor

```

```

..
#         class_names = [ 'machine_status'],
#         rounded = True, proportion = False,
#         precision = 2, filled = True)

# validacion cruzada
# https://jamesrledoux.com/code/k\_fold\_cross\_validation

from sklearn.model_selection import cross_validate

start1 = time.time()
model = RandomForestClassifier(random_state=1)
cv = cross_validate(model, X, y, cv=10)
print(cv['test_score'])
print(cv['test_score'].mean())
stop1 = time.time()
print(f"Training time: {stop1 - start1}s")

[0.99464477 0.99964298 1.          1.          1.          1.
 1.          1.          1.          0.85642857]
0.9850716325802009
Training time: 30.822885751724243s

https://stackoverflow.com/questions/20662023/save-python-random-forest-model-to-file

from sklearn.metrics import classification_report, confusion_matrix
import matplotlib.pyplot as plt
from sklearn.metrics import plot_confusion_matrix

start1 = time.time()
#model = RandomForestClassifier(random_state=1)
model = RandomForestClassifier(n_estimators=25)

cv = cross_validate(model, X, y, cv=10)
print(confusion_matrix(y_test,y_pred))
print(cv['test_score'])
print(cv['test_score'].mean())
stop1 = time.time()
print(f"Training time: {stop1 - start1}s")

print(confusion_matrix(y_test,y_pred))
print(classification_report(y_test,y_pred))

#plot_confusion_matrix(clf, X_test, y_test)
# plot_confusion_matrix(clf, X_test, y_test)
plot_confusion_matrix(clf, X_test, y_test, cmap=plt.cm.Blues)

plt.show()

```


```
[[2789    0]
 [    1 2811]]
[0.99107462 0.99964298 1.          1.          0.99892857 1.
 1.          1.          1.          0.85535714]
```

```
0.9845003315142551
```

```
Training time: 7.9669787883758545s
```

```
[[2789    0]
 [    1 2811]]
```

		precision	recall	f1-score	support
	0	1.00	1.00	1.00	2789
	1	1.00	1.00	1.00	2812
accuracy				1.00	5601
macro avg		1.00	1.00	1.00	5601
weighted avg		1.00	1.00	1.00	5601



```
# version with multi scoring
from sklearn.metrics import classification_report, confusion_matrix
import matplotlib.pyplot as plt
from sklearn.metrics import plot_confusion_matrix

start1 = time.time()
#model = RandomForestClassifier(random_state=1)
model = RandomForestClassifier(n_estimators=25)

cv = cross_validate(model, X, y, cv=10)
#recall_score=cross_validation.cross_val_score(clf, X,y, cv=10, scoring ='recall')
#recall_score=cross_val_score(model, X,y, cv=10, scoring ='recall')
f1=cross_validate(model, X,y, cv=10, scoring ='f1')
recall_score=cross_validate(model, X,y, cv=10, scoring ='recall')
pre_score=cross_validate(model, X,y, cv=10, scoring ='precision_macro')
print(confusion_matrix(y_test,y_pred))
print(f"precision_macro_score:")
print(pre_score['test_score'])
print(pre_score['test_score'].mean())
print(f"test_score:")
print(cv['test_score'])
```



```
print(cv['test_score'].mean())
print(f"recall:")
print(recall_score['test_score'])
print(recall_score['test_score'].mean())
print(f"f1score:")
print(f1['test_score'])
print(f1['test_score'].mean())
stop1 = time.time()
print(f"Training time: {stop1 - start1}s")

print(confusion_matrix(y_test,y_pred))
print(classification_report(y_test,y_pred))

#plot_confusion_matrix(clf, X_test, y_test)
# plot_confusion_matrix(clf, X_test, y_test)
plot_confusion_matrix(clf, X_test, y_test, cmap=plt.cm.Blues)

plt.show()
```

```

[[2789    0]
 [   1 2811]]
precision_macro_score:
[0.99680851 0.99964311 1.          0.99751244 1.          1.
 1.          1.          1.          0.92250945]
0.9916473508329371
test_score:
[0.99714388 0.99964298 1.          1.          1.          1.
 1.          1.          1.          0.79071429]
0.97075011117510215

```

```
# version with multi scoring mejorada
```

```
from sklearn.metrics import classification_report, confusion_matrix
```

```
import matplotlib.pyplot as plt
```

```
from sklearn.metrics import plot_confusion_matrix
```

```
start1 = time.time()
```

```
#model = RandomForestClassifier(random_state=1)
```

```
model = RandomForestClassifier(n_estimators=25)
```

```
#GH
```

```
model.fit(X_train,y_train)
```

```
y_pred=model.predict(X_test)
```

```
#GH
```

```
cv = cross_validate(model, X, y, cv=10)
```

```
#recall_score=cross_validation.cross_val_score(clf, X,y, cv=10, scoring ='recall')
```

```
#recall_score=cross_val_score(model, X,y, cv=10, scoring ='recall')
```

```
#scoring = ['neg_mean_absolute_error','r2']
```

```
scores=cross_validate(model, X,y, cv=10, scoring = ['accuracy','f1','recall','precision'],ret
```

```
#recall_score=cross_validate(model, X,y, cv=10, scoring ='recall')
```

```
#pre_score=cross_validate(model, X,y, cv=10, scoring ='precision_macro')
```

```
print(confusion_matrix(y_test,y_pred))
```

```
print(f"multi_metric_scores:")
```

```
#print(scores['test_score'])
```

```
print(scores)
```

```
#print(scores['test_score'].mean())
```

```
#print(scores.mean())
```

```
#print(f"precision_macro_score:")
```

```
#print(pre_score['test_score'])
```

```
#print(pre_score['test_score'].mean())
```

```
#print(f"test_score:")
```

```
#print(cv['test_score'])
```

```
#print(cv['test_score'].mean())
```

```
#print(f"recall:")
```

```
#print(recall_score['test_score'])
```

```
#print(recall_score['test_score'].mean())

#print(f"f1score:")
#print(f1['test_score'])
#print(f1['test_score'].mean())

stop1 = time.time()
print(f"Training time: {stop1 - start1}s")

print(confusion_matrix(y_test,y_pred))
print(classification_report(y_test,y_pred))

#plot_confusion_matrix(clf, X_test, y_test)
# plot_confusion_matrix(clf, X_test, y_test)
plot_confusion_matrix(clf, X_test, y_test, cmap=plt.cm.Blues)

plt.show()

gh4 = scores.get("test_accuracy")

print(f"accuracy:")
print(gh4)
print(gh4.mean())

gh3 = scores.get("test_precision")

print(f"precision:")
print(gh3)
print(gh3.mean())

gh = scores.get("test_recall")

print(f"recall:")
print(gh)
print(gh.mean())

gh2 = scores.get("test_f1")

print(f"f1:")
print(gh2)
print(gh2.mean())

CM = confusion_matrix(y_test, y_pred)
print(f"-----")
print(f"matriz de confusion:")
TN = CM[0][0]
FN = CM[1][0]
TP = CM[1][1]
FP = CM[0][1]
print(f"TN={TN}, FP={FP} ")
print(f"FN={FN}, TP={TP} ")
```

```
print(f"-----")
print(f"matriz de confusion %:")
total1=(TN+TP+FN+FP)

print(f"TN={100*TN/total1}, FP={100*FP/total1} ")
print(f"FN={100*FN/total1}, TP={100*TP/total1} ")

print(f"-----")
acc1=(TN+TP)/(TN+TP+FN+FP)
print(f"accuracy1={acc1}")

print(f"-----")
re1=(TP)/(TP+FN)
print(f"reca1={re1}")

print(f"-----")
pre1=(TP)/(TP+FP)
print(f"pre1={pre1}")

print(f"-----")
f1s1=(2*pre1*re1)/(pre1+re1)
print(f"f1score={f1s1}")
```


```
[[2789    0]
 [    1 2811]]
multi_metric_scores:
{'fit_time': array([0.68429852, 0.75059485, 0.79937196, 0.81243873, 0.76197147,
 0.78753948, 0.76873612, 0.80026364, 0.79576921, 0.83706284]), 'score_time': array(
0.01073599, 0.01081562, 0.01152539, 0.01071858, 0.01142955]), 'test_accuracy': ar
1.          , 1.          , 1.          , 1.          , 0.79464286]), 'test_f1': array([0
1.          , 1.          , 1.          , 1.          , 0.82952861]), 'test_recall': arra
1.          , 1.          , 1.          , 1.          , 0.99928571]), 'test_precision': a
1.          , 1.          , 1.          , 1.          , 0.70907248]}}
```

Training time: 16.54348850250244s

```
[[2789    0]
 [    1 2811]]
      precision    recall  f1-score   support

     0       1.00      1.00      1.00     2789
     1       1.00      1.00      1.00     2812

 accuracy          1.00          1.00          1.00          5601
 macro avg       1.00      1.00      1.00          5601
weighted avg       1.00      1.00      1.00          5601
```



accuracy:

```
import joblib
from sklearn.ensemble import RandomForestClassifier
# create RF
clf = RandomForestClassifier(n_estimators=100, random_state=42)
# fit
clf.fit(X_train, y_train)
# save
joblib.dump(clf, "my_random_forest.joblib")

# load
loaded_rf = joblib.load("my_random_forest.joblib")

# predicciones
#clf.predict([[0.0,53.55902,0,0,0,0,0,0,0,0,0,0,0,0,0,0,0,0,0,0]])
```

```

#predicciones
loaded_rf.predict([[0.0,53.55902,0,0,0,0,0,0,0,0,0,0,0,0,0,0,0,0,0,0]])

    array([1])
    -----

# 1 es recovering
loaded_rf.predict([[0.0,0,0,0,0,0,0,0,0,0,0,0,0,0,0,0,0,0,0,0]])

    array([1])
    ~~~~~

# 0 es recovering
loaded_rf.predict([[2.465394,47.092009999999995,53.2118,46.310759999999995,634375,47.52422,41

    array([1])

# 2 es broken
loaded_rf.predict([[2.258796,47.26563,52.73437,43.4461784362793,200.11573791503898,43.62322,4

    array([1])

import pandas as pd
#feature_imp = pd.Series(clf.feature_importances_,index=iris.feature_names).sort_values(ascen
#feature_imp = pd.Series(clf.feature_importances_,index=X.columns[1:8]).sort_values(ascending
feature_imp = pd.Series(clf.feature_importances_,index=X.columns[0:24]).sort_values(ascending
print(feature_imp)

#Visualización
import matplotlib.pyplot as plt
import seaborn as sns
%matplotlib inline
# Creating a bar plot
sns.barplot(x=feature_imp, y=feature_imp.index)
# Add labels to your graph
#plt.xlabel('Feature Importance Score')
plt.xlabel('Score de Caracteristicas importantes')
#plt.ylabel('Features')
plt.ylabel('Caracteristicas')
#plt.title("Visualizing Important Features")
plt.title("Visualización de carasteristicas importantes")
plt.legend()
plt.show()

#plt.savefig('destination_path.eps', format='eps' , dpi=1000)

plt.savefig('myimage.svg', format='svg', dpi=1200)

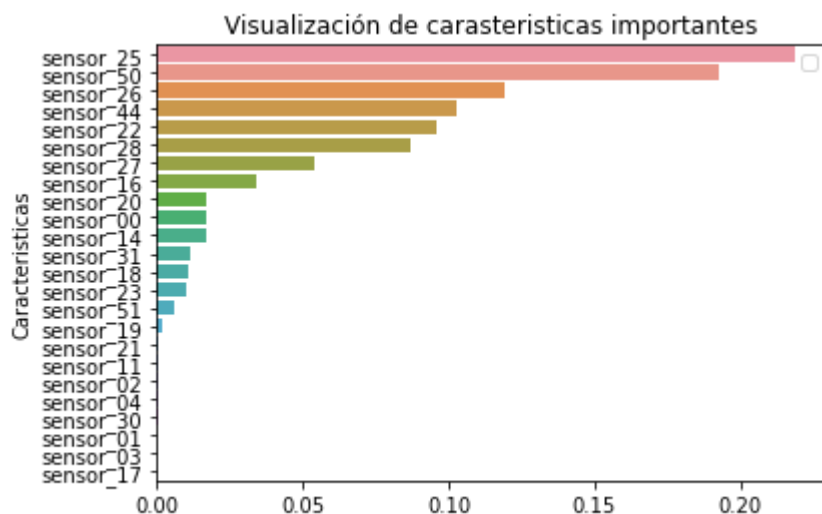
```

```

sensor_25    0.218585577489896
sensor_50    0.192494505492431
sensor_26    0.118933369964408
sensor_44    0.103024605453617
sensor_22    0.095838605122650
sensor_28    0.086903813644770
sensor_27    0.054355455121461
sensor_16    0.034016238185235
sensor_20    0.017186116553708
sensor_00    0.017118197083998
sensor_14    0.017109034806036
sensor_31    0.011367880294808
sensor_18    0.010654964736604
sensor_23    0.010360987502906
sensor_51    0.006311624990248
sensor_19    0.002119835128739
sensor_21    0.000733754684364
sensor_11    0.000626502174069
sensor_02    0.000602177105466
sensor_04    0.000544201099041
sensor_30    0.000456629591572
sensor_01    0.000279868153119
sensor_03    0.000258636462968
sensor_17    0.000117419157886
dtype: float64

```

No handles with labels found to put in legend.



<https://towardsdatascience.com/how-to-visualize-a-decision-tree-from-a-random-forest-in-pyt>